

```
In [1]: %matplotlib inline

import matplotlib.pyplot as plt
from keras.utils import np_utils
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization

import seaborn as sns
import numpy as np
import keras
```

Using TensorFlow backend.

```
In [3]: (X_train, y_train), (X_test, y_test) = mnist.load_data()

img_rows, img_cols = 28, 28
```

```
In [4]: #Let's print the shape of training and test data

print("Training data shape: ", X_train.shape)
print("Test data shape", X_test.shape)

print("Training label shape: ", y_train.shape)
print("Test label shape: ", y_test.shape)
print("First 5 training labels: ", y_train[:5])

Training data shape: (60000, 28, 28)
Test data shape (10000, 28, 28)
Training label shape: (60000,)
Test label shape: (10000,)
First 5 training labels: [5 0 4 1 9]
```

```
In [5]: num_classes = 10
batch_size = 128
nb_epoch = 10
```

```
In [7]: from keras import backend as K

if K.image_data_format() == 'channels_first':
    X_train = X_train.reshape(X_train.shape[0], 1, img_rows, img_cols)
    X_test = X_test.reshape(X_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    X_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1)
    X_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
```

```
In [8]: # Each cell of above matrix is having a value between 0-255
# before applying machine learning algorithms, Let's normalize the data #  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train /= 255
X_test /= 255
print('X_train shape:', X_train.shape)
print('No. of train samples: ', X_train.shape[0])
print('No. of test samples: ', X_test.shape[0])

X_train shape: (60000, 28, 28, 1)
No. of train samples: 60000
No. of test samples: 10000
```

```
In [9]: # Let's convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

4 layer CNN Model with 3*3 kernel size

```
In [10]: model1=Sequential() # Initializing the model

# First ConvNet
model1.add(Conv2D(32,kernel_size=(3,3),
                  activation='relu',
                  input_shape=input_shape))

model1.add(Conv2D(64,kernel_size=(3,3),
                  activation='relu'))

model1.add(MaxPooling2D(pool_size=(2,2)))
model1.add(Dropout(0.2))

model1.add(Conv2D(128,kernel_size=(3,3),
                  activation='relu'))

model1.add(BatchNormalization())

#maxpooling by (2,2) ,dropout,flattening
model1.add(MaxPooling2D(pool_size=(2,2)))

model1.add(Dropout(0.5))

model1.add(Flatten())

#hidden_layer
model1.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model1.add(Dropout(0.5))

model1.add(Dense(num_classes,activation='softmax'))

print(model1.summary())
```

WARNING:tensorflow:From C:\Users\pratyush.acharya\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\ops\resource_variable_ops.py:435: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
conv2d_3 (Conv2D)	(None, 10, 10, 128)	73856
batch_normalization_1 (Batch Normalization)	(None, 10, 10, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 128)	0
dropout_2 (Dropout)	(None, 5, 5, 128)	0
flatten_1 (Flatten)	(None, 3200)	0
dense_1 (Dense)	(None, 256)	819456
dropout_3 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2570
Total params: 915,210		
Trainable params: 914,954		
Non-trainable params: 256		
None		

```
In [11]: model1.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [12]: # Training the model

history = model1.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))

WARNING:tensorflow:From C:\Users\pratyush.acharya\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.cast instead.
Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 527s 9ms/step - loss: 0.2391 - accuracy: 0.9286 - val_loss: 0.1245 - val_accuracy: 0.9855
Epoch 2/10
60000/60000 [=====] - 537s 9ms/step - loss: 0.0807 - accuracy: 0.9751 - val_loss: 0.0440 - val_accuracy: 0.9860
Epoch 3/10
60000/60000 [=====] - 535s 9ms/step - loss: 0.0623 - accuracy: 0.9816 - val_loss: 0.0283 - val_accuracy: 0.9910
Epoch 4/10
60000/60000 [=====] - 551s 9ms/step - loss: 0.0544 - accuracy: 0.9839 - val_loss: 0.0309 - val_accuracy: 0.9907
Epoch 5/10
60000/60000 [=====] - 550s 9ms/step - loss: 0.0509 - accuracy: 0.9850 - val_loss: 0.0281 - val_accuracy: 0.9919
Epoch 6/10
60000/60000 [=====] - 539s 9ms/step - loss: 0.0459 - accuracy: 0.9860 - val_loss: 0.0221 - val_accuracy: 0.9921
Epoch 7/10
60000/60000 [=====] - 464s 8ms/step - loss: 0.0413 - accuracy: 0.9874 - val_loss: 0.0227 - val_accuracy: 0.9933
Epoch 8/10
60000/60000 [=====] - 376s 6ms/step - loss: 0.0427 - accuracy: 0.9875 - val_loss: 0.0301 - val_accuracy: 0.9925
Epoch 9/10
60000/60000 [=====] - 342s 6ms/step - loss: 0.0409 - accuracy: 0.9883 - val_loss: 0.0371 - val_accuracy: 0.9908
Epoch 10/10
60000/60000 [=====] - 339s 6ms/step - loss: 0.0373 - accuracy: 0.9893 - val_loss: 0.0220 - val_accuracy: 0.9930
```

```
In [13]: score = model1.evaluate(X_test, y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])


10000/10000 [=====] - 14s 1ms/step
Test loss: 0.022041849251561642
Test accuracy: 0.9929999709129333
```

```
In [14]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



Changing the dropout rate to 0.8

```
In [15]: model1=Sequential() # Initializing the model

# First ConvNet
model1.add(Conv2D(32,kernel_size=(3,3),
                  activation='relu',
                  input_shape=input_shape))

model1.add(Conv2D(64,kernel_size=(3,3),
                  activation='relu'))

model1.add(MaxPooling2D(pool_size=(2,2)))
model1.add(Dropout(0.8))

model1.add(Conv2D(128,kernel_size=(3,3),
                  activation='relu'))

model1.add(BatchNormalization())

#maxpooling by (2,2) ,dropout,flattening
model1.add(MaxPooling2D(pool_size=(2,2)))

model1.add(Dropout(0.8))

model1.add(Flatten())

#hidden_layer
model1.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model1.add(Dropout(0.8))

model1.add(Dense(num_classes,activation='softmax'))

print(model1.summary())
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 26, 26, 32)	320
conv2d_5 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_4 (Dropout)	(None, 12, 12, 64)	0
conv2d_6 (Conv2D)	(None, 10, 10, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 10, 10, 128)	512
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 128)	0
dropout_5 (Dropout)	(None, 5, 5, 128)	0
flatten_2 (Flatten)	(None, 3200)	0
dense_3 (Dense)	(None, 256)	819456
dropout_6 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2570
=====		
Total params: 915,210		
Trainable params: 914,954		
Non-trainable params: 256		
None		

```
In [16]: model1.compile(optimizer='adam',
                       loss='categorical_crossentropy',
                       metrics=['accuracy'])
```

In [17]: *# Training the model*

```
history = model1.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10
60000/60000 [=====] - 354s 6ms/step - loss: 1.8717 - accuracy: 0.4944 - val_loss: 1.2311 - val_accuracy: 0.5943

Epoch 2/10
60000/60000 [=====] - 339s 6ms/step - loss: 0.7126 - accuracy: 0.7703 - val_loss: 0.3755 - val_accuracy: 0.8838

Epoch 3/10
60000/60000 [=====] - 337s 6ms/step - loss: 0.5028 - accuracy: 0.8438 - val_loss: 0.1594 - val_accuracy: 0.9600

Epoch 4/10
60000/60000 [=====] - 407s 7ms/step - loss: 0.3878 - accuracy: 0.8838 - val_loss: 0.3313 - val_accuracy: 0.9028

Epoch 5/10
60000/60000 [=====] - 420s 7ms/step - loss: 0.3306 - accuracy: 0.9019 - val_loss: 0.1561 - val_accuracy: 0.9538

Epoch 6/10
60000/60000 [=====] - 403s 7ms/step - loss: 0.2974 - accuracy: 0.9145 - val_loss: 0.1993 - val_accuracy: 0.9451

Epoch 7/10
60000/60000 [=====] - 385s 6ms/step - loss: 0.2826 - accuracy: 0.9189 - val_loss: 0.1173 - val_accuracy: 0.9664

Epoch 8/10
60000/60000 [=====] - 371s 6ms/step - loss: 0.2609 - accuracy: 0.9247 - val_loss: 0.1695 - val_accuracy: 0.9550

Epoch 9/10
60000/60000 [=====] - 359s 6ms/step - loss: 0.2491 - accuracy: 0.9276 - val_loss: 0.1613 - val_accuracy: 0.9575

Epoch 10/10
60000/60000 [=====] - 340s 6ms/step - loss: 0.2467 - accuracy: 0.9296 - val_loss: 0.2272 - val_accuracy: 0.9399

In [18]: `score = model1.evaluate(X_test, y_test, verbose=1)`
`print('Test loss:', score[0])`
`print('Test accuracy:', score[1])`


10000/10000 [=====] - 16s 2ms/step
Test loss: 0.22723832533806562
Test accuracy: 0.9398999810218811

In [20]: `fig,ax = plt.subplots(1,1)`
`ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')`

List of epoch numbers
`list_of_epoch = list(range(1,nb_epoch+1))`

`train_loss = history.history['loss']`
`val_loss = history.history['val_loss']`

`ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")`
`ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")`
`plt.legend()`
`plt.grid()`
`plt.show()`



Changing the dropout rate to 0.2

```
In [21]: model1=Sequential() # Initializing the model

# First ConvNet
model1.add(Conv2D(32,kernel_size=(3,3),
                  activation='relu',
                  input_shape=input_shape))

model1.add(Conv2D(64,kernel_size=(3,3),
                  activation='relu'))

model1.add(MaxPooling2D(pool_size=(2,2)))
model1.add(Dropout(0.2))

model1.add(Conv2D(128,kernel_size=(3,3),
                  activation='relu'))

model1.add(BatchNormalization())

#maxpooling by (2,2) ,dropout,flattening
model1.add(MaxPooling2D(pool_size=(2,2)))

model1.add(Dropout(0.2))

model1.add(Flatten())

#hidden_layer
model1.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model1.add(Dropout(0.2))

model1.add(Dense(num_classes,activation='softmax'))

print(model1.summary())
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 26, 26, 32)	320
conv2d_8 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_7 (Dropout)	(None, 12, 12, 64)	0
conv2d_9 (Conv2D)	(None, 10, 10, 128)	73856
batch_normalization_3 (Batch Normalization)	(None, 10, 10, 128)	512
max_pooling2d_6 (MaxPooling2D)	(None, 5, 5, 128)	0
dropout_8 (Dropout)	(None, 5, 5, 128)	0
flatten_3 (Flatten)	(None, 3200)	0
dense_5 (Dense)	(None, 256)	819456
dropout_9 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 10)	2570
=====		
Total params: 915,210		
Trainable params: 914,954		
Non-trainable params: 256		
None		

```
In [22]: model1.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

In [23]: *# Training the model*

```
history = model1.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10
60000/60000 [=====] - 442s 7ms/step - loss: 0.1491 - accuracy: 0.9553 - val_loss: 0.0955 - val_accuracy: 0.9895

Epoch 2/10
60000/60000 [=====] - 387s 6ms/step - loss: 0.0493 - accuracy: 0.9851 - val_loss: 0.0546 - val_accuracy: 0.9840

Epoch 3/10
60000/60000 [=====] - 416s 7ms/step - loss: 0.0363 - accuracy: 0.9891 - val_loss: 0.0264 - val_accuracy: 0.9914

Epoch 4/10
60000/60000 [=====] - 394s 7ms/step - loss: 0.0290 - accuracy: 0.9911 - val_loss: 0.0365 - val_accuracy: 0.9894

Epoch 5/10
60000/60000 [=====] - 449s 7ms/step - loss: 0.0246 - accuracy: 0.9923 - val_loss: 0.0298 - val_accuracy: 0.9920

Epoch 6/10
60000/60000 [=====] - 410s 7ms/step - loss: 0.0219 - accuracy: 0.9933 - val_loss: 0.0280 - val_accuracy: 0.9917

Epoch 7/10
60000/60000 [=====] - 366s 6ms/step - loss: 0.0214 - accuracy: 0.9931 - val_loss: 0.0283 - val_accuracy: 0.9921

Epoch 8/10
60000/60000 [=====] - 376s 6ms/step - loss: 0.0190 - accuracy: 0.9940 - val_loss: 0.0383 - val_accuracy: 0.9910

Epoch 9/10
60000/60000 [=====] - 372s 6ms/step - loss: 0.0170 - accuracy: 0.9948 - val_loss: 0.0303 - val_accuracy: 0.9919

Epoch 10/10
60000/60000 [=====] - 397s 7ms/step - loss: 0.0181 - accuracy: 0.9945 - val_loss: 0.0351 - val_accuracy: 0.9921

In [24]: `score = model1.evaluate(X_test, y_test, verbose=1)`
`print('Test loss:', score[0])`
`print('Test accuracy:', score[1])`


```
10000/10000 [=====] - 16s 2ms/step
Test loss: 0.03512107918624624
Test accuracy: 0.992100003814697
```

In [25]: `fig,ax = plt.subplots(1,1)`
`ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')`

List of epoch numbers
`list_of_epoch = list(range(1,nb_epoch+1))`

`train_loss = history.history['loss']`
`val_loss = history.history['val_loss']`

`ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")`
`ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")`
`plt.legend()`
`plt.grid()`
`plt.show();`



4 Convolution layers with 5*5 kernel size

```
In [27]: model2=Sequential() # Initializing the model

# First ConvNet
model2.add(Conv2D(32,kernel_size=(5,5),
                  activation='relu',
                  input_shape=input_shape))

model2.add(Conv2D(64,kernel_size=(5,5),
                  activation='relu'))

model2.add(MaxPooling2D(pool_size=(2,2)))
model2.add(Dropout(0.2))

model2.add(Conv2D(128,kernel_size=(5,5),
                  activation='relu'))

model2.add(BatchNormalization())

model2.add(MaxPooling2D(pool_size=(2,2)))
model2.add(Dropout(0.5))

model2.add(Flatten())

#hidden layer
model2.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model2.add(Dropout(0.5))

model2.add(Dense(num_classes,activation='softmax'))

print(model2.summary())
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 24, 24, 32)	832
conv2d_11 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_7 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_10 (Dropout)	(None, 10, 10, 64)	0
conv2d_12 (Conv2D)	(None, 6, 6, 128)	204928
batch_normalization_4 (Batch Normalization)	(None, 6, 6, 128)	512
max_pooling2d_8 (MaxPooling2D)	(None, 3, 3, 128)	0
dropout_11 (Dropout)	(None, 3, 3, 128)	0
flatten_4 (Flatten)	(None, 1152)	0
dense_7 (Dense)	(None, 256)	295168
dropout_12 (Dropout)	(None, 256)	0
dense_8 (Dense)	(None, 10)	2570
Total params: 555,274		
Trainable params: 555,018		
Non-trainable params: 256		
None		

```
In [28]: model2.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```



```
In [29]: # Training the model
history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 359s 6ms/step - loss: 0.1954 - accuracy: 0.9391 - val_loss: 0.0761 - val_accuracy: 0.9809
Epoch 2/10
60000/60000 [=====] - 361s 6ms/step - loss: 0.0752 - accuracy: 0.9781 - val_loss: 0.0286 - val_accuracy: 0.9913
Epoch 3/10
60000/60000 [=====] - 341s 6ms/step - loss: 0.0565 - accuracy: 0.9831 - val_loss: 0.0257 - val_accuracy: 0.9912
Epoch 4/10
60000/60000 [=====] - 465s 8ms/step - loss: 0.0488 - accuracy: 0.9856 - val_loss: 0.0261 - val_accuracy: 0.9906
Epoch 5/10
60000/60000 [=====] - 573s 10ms/step - loss: 0.0440 - accuracy: 0.9865 - val_loss: 0.0332 - val_accuracy: 0.9911
Epoch 6/10
60000/60000 [=====] - 540s 9ms/step - loss: 0.0422 - accuracy: 0.9878 - val_loss: 0.0309 - val_accuracy: 0.9924
Epoch 7/10
60000/60000 [=====] - 567s 9ms/step - loss: 0.0398 - accuracy: 0.9884 - val_loss: 0.0296 - val_accuracy: 0.9911
Epoch 8/10
60000/60000 [=====] - 572s 10ms/step - loss: 0.0374 - accuracy: 0.9888 - val_loss: 0.0222 - val_accuracy: 0.9936
Epoch 9/10
60000/60000 [=====] - 561s 9ms/step - loss: 0.0321 - accuracy: 0.9904 - val_loss: 0.0240 - val_accuracy: 0.9938
Epoch 10/10
60000/60000 [=====] - 547s 9ms/step - loss: 0.0321 - accuracy: 0.9905 - val_loss: 0.0210 - val_accuracy: 0.9943
```

```
In [30]: score = model2.evaluate(X_test, y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])


10000/10000 [=====] - 27s 3ms/step
Test loss: 0.020990679916361204
Test accuracy: 0.9943000078201294
```

```
In [31]: fig, ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1, nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show()
```



Changing the dropout rate to 0.8

```
In [32]: model2=Sequential() # Initializing the model

# First ConvNet
model2.add(Conv2D(32, kernel_size=(5,5),
                  activation='relu',
                  input_shape=input_shape))

model2.add(Conv2D(64, kernel_size=(5,5),
                  activation='relu'))

model2.add(MaxPooling2D(pool_size=(2,2)))
model2.add(Dropout(0.8))

model2.add(Conv2D(128, kernel_size=(5,5),
                  activation='relu'))

model2.add(BatchNormalization())

model2.add(MaxPooling2D(pool_size=(2,2)))
model2.add(Dropout(0.8))

model2.add(Flatten())

#hidden layer
model2.add(Dense(256, activation='relu', kernel_initializer='random_uniform'))

model2.add(Dropout(0.8))

model2.add(Dense(num_classes, activation='softmax'))

print(model2.summary())
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 24, 24, 32)	832
conv2d_14 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_9 (MaxPooling2D)	(None, 10, 10, 64)	0
dropout_13 (Dropout)	(None, 10, 10, 64)	0
conv2d_15 (Conv2D)	(None, 6, 6, 128)	204928
batch_normalization_5 (Batch Normalization)	(None, 6, 6, 128)	512
max_pooling2d_10 (MaxPooling2D)	(None, 3, 3, 128)	0
dropout_14 (Dropout)	(None, 3, 3, 128)	0
flatten_5 (Flatten)	(None, 1152)	0
dense_9 (Dense)	(None, 256)	295168
dropout_15 (Dropout)	(None, 256)	0
dense_10 (Dense)	(None, 10)	2570
Total params: 555,274		
Trainable params: 555,018		
Non-trainable params: 256		
None		

```
In [33]: model2.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

In [34]: *# Training the model*

```
history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10
60000/60000 [=====] - 548s 9ms/step - loss: 1.6143 - accuracy: 0.4881 - val_loss: 0.7136 - val_accuracy: 0.8021

Epoch 2/10
60000/60000 [=====] - 540s 9ms/step - loss: 0.4909 - accuracy: 0.8488 - val_loss: 0.1857 - val_accuracy: 0.9520

Epoch 3/10
60000/60000 [=====] - 548s 9ms/step - loss: 0.3236 - accuracy: 0.9045 - val_loss: 0.0913 - val_accuracy: 0.9740

Epoch 4/10
60000/60000 [=====] - 530s 9ms/step - loss: 0.2664 - accuracy: 0.9242 - val_loss: 0.1157 - val_accuracy: 0.9725

Epoch 5/10
60000/60000 [=====] - 533s 9ms/step - loss: 0.2406 - accuracy: 0.9318 - val_loss: 0.0731 - val_accuracy: 0.9802

Epoch 6/10
60000/60000 [=====] - 532s 9ms/step - loss: 0.2206 - accuracy: 0.9374 - val_loss: 0.1145 - val_accuracy: 0.9733

Epoch 7/10
60000/60000 [=====] - 545s 9ms/step - loss: 0.2124 - accuracy: 0.9405 - val_loss: 0.1015 - val_accuracy: 0.9770

Epoch 8/10
60000/60000 [=====] - 532s 9ms/step - loss: 0.2043 - accuracy: 0.9439 - val_loss: 0.1037 - val_accuracy: 0.9768

Epoch 9/10
60000/60000 [=====] - 529s 9ms/step - loss: 0.1905 - accuracy: 0.9470 - val_loss: 0.0792 - val_accuracy: 0.9814

Epoch 10/10
60000/60000 [=====] - 543s 9ms/step - loss: 0.1791 - accuracy: 0.9492 - val_loss: 0.0698 - val_accuracy: 0.9834

In [35]:

```
score = model2.evaluate(X_test, y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
10000/10000 [=====] - 24s 2ms/step
Test loss: 0.06979307891329391
Test accuracy: 0.9833999872207642
```


In [36]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show()
```



Changing the dropout to 0.2

```
In [37]: model2=Sequential() # Initializing the model

# First ConvNet
model2.add(Conv2D(32,kernel_size=(5,5),
                  activation='relu',
                  input_shape=input_shape))

model2.add(Conv2D(64,kernel_size=(5,5),
                  activation='relu'))

model2.add(MaxPooling2D(pool_size=(2,2)))
model2.add(Dropout(0.2))

model2.add(Conv2D(128,kernel_size=(5,5),
                  activation='relu'))

model2.add(BatchNormalization())

model2.add(MaxPooling2D(pool_size=(2,2)))
model2.add(Dropout(0.2))

model2.add(Flatten())

#hidden layer
model2.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model2.add(Dropout(0.2))

model2.add(Dense(num_classes,activation='softmax'))

print(model2.summary())
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
conv2d_16 (Conv2D)	(None, 24, 24, 32)	832
conv2d_17 (Conv2D)	(None, 20, 20, 64)	51264
max_pooling2d_11 (MaxPooling)	(None, 10, 10, 64)	0
dropout_16 (Dropout)	(None, 10, 10, 64)	0
conv2d_18 (Conv2D)	(None, 6, 6, 128)	204928
batch_normalization_6 (Batch Normalization)	(None, 6, 6, 128)	512
max_pooling2d_12 (MaxPooling)	(None, 3, 3, 128)	0
dropout_17 (Dropout)	(None, 3, 3, 128)	0
flatten_6 (Flatten)	(None, 1152)	0
dense_11 (Dense)	(None, 256)	295168
dropout_18 (Dropout)	(None, 256)	0
dense_12 (Dense)	(None, 10)	2570
Total params: 555,274		
Trainable params: 555,018		
Non-trainable params: 256		
None		

```
In [38]: model2.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

In [39]: *# Training the model*

```
history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10
60000/60000 [=====] - 534s 9ms/step - loss: 0.1239 - accuracy: 0.9619 - val_loss: 0.0679 - val_accuracy: 0.9805

Epoch 2/10
60000/60000 [=====] - 532s 9ms/step - loss: 0.0471 - accuracy: 0.9856 - val_loss: 0.0539 - val_accuracy: 0.9820

Epoch 3/10
60000/60000 [=====] - 444s 7ms/step - loss: 0.0370 - accuracy: 0.9887 - val_loss: 0.0362 - val_accuracy: 0.9888

Epoch 4/10
60000/60000 [=====] - 332s 6ms/step - loss: 0.0287 - accuracy: 0.9913 - val_loss: 0.0402 - val_accuracy: 0.9890

Epoch 5/10
60000/60000 [=====] - 330s 6ms/step - loss: 0.0256 - accuracy: 0.9924 - val_loss: 0.0268 - val_accuracy: 0.9916

Epoch 6/10
60000/60000 [=====] - 329s 5ms/step - loss: 0.0210 - accuracy: 0.9938 - val_loss: 0.0273 - val_accuracy: 0.9916

Epoch 7/10
60000/60000 [=====] - 338s 6ms/step - loss: 0.0189 - accuracy: 0.9943 - val_loss: 0.0347 - val_accuracy: 0.9918

Epoch 8/10
60000/60000 [=====] - 359s 6ms/step - loss: 0.0191 - accuracy: 0.9940 - val_loss: 0.0238 - val_accuracy: 0.9940

Epoch 9/10
60000/60000 [=====] - 339s 6ms/step - loss: 0.0148 - accuracy: 0.9951 - val_loss: 0.0271 - val_accuracy: 0.9939

Epoch 10/10
60000/60000 [=====] - 331s 6ms/step - loss: 0.0165 - accuracy: 0.9946 - val_loss: 0.0278 - val_accuracy: 0.9930

In [40]: `score = model2.evaluate(X_test, y_test, verbose=1)`
`print('Test loss:', score[0])`
`print('Test accuracy:', score[1])`


```
10000/10000 [=====] - 19s 2ms/step
Test loss: 0.027773732883542696
Test accuracy: 0.9929999709129333
```

In [41]: `fig,ax = plt.subplots(1,1)`
`ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')`

List of epoch numbers
`list_of_epoch = list(range(1,nb_epoch+1))`

`train_loss = history.history['loss']`
`val_loss = history.history['val_loss']`

`ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")`
`ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")`
`plt.legend()`
`plt.grid()`
`plt.show()`



3 layer convolution layer with 3*3 kernel size

```
In [42]: model3=Sequential() # Initializing the model

# First ConvNet
model3.add(Conv2D(32,kernel_size=(3,3),
                  activation='relu',
                  input_shape=input_shape))
model3.add(MaxPooling2D(pool_size=(2,2)))

model3.add(Conv2D(64,kernel_size=(3,3),
                  activation='relu',padding='valid'))

model3.add(Dropout(0.5))

#hidden_layer
model3.add(Flatten())

model3.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))
model3.add(Dense(num_classes,activation='softmax'))

print(model3.summary())
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
conv2d_19 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_13 (MaxPooling)	(None, 13, 13, 32)	0
conv2d_20 (Conv2D)	(None, 11, 11, 64)	18496
dropout_19 (Dropout)	(None, 11, 11, 64)	0
flatten_7 (Flatten)	(None, 7744)	0
dense_13 (Dense)	(None, 256)	1982720
dense_14 (Dense)	(None, 10)	2570
=====		
Total params: 2,004,106		
Trainable params: 2,004,106		
Non-trainable params: 0		
None		

```
In [43]: model3.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

```
In [44]: # Training the model

history = model3.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 92s 2ms/step - loss: 0.1748 - accuracy: 0.9464 - val_loss: 0.0502 - val_accuracy: 0.9840
Epoch 2/10
60000/60000 [=====] - 89s 1ms/step - loss: 0.0566 - accuracy: 0.9824 - val_loss: 0.0368 - val_accuracy: 0.9875
Epoch 3/10
60000/60000 [=====] - 88s 1ms/step - loss: 0.0390 - accuracy: 0.9876 - val_loss: 0.0301 - val_accuracy: 0.9900
Epoch 4/10
60000/60000 [=====] - 87s 1ms/step - loss: 0.0302 - accuracy: 0.9906 - val_loss: 0.0304 - val_accuracy: 0.9900
Epoch 5/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0238 - accuracy: 0.9922 - val_loss: 0.0262 - val_accuracy: 0.9915
Epoch 6/10
60000/60000 [=====] - 86s 1ms/step - loss: 0.0197 - accuracy: 0.9933 - val_loss: 0.0260 - val_accuracy: 0.9918
Epoch 7/10
60000/60000 [=====] - 89s 1ms/step - loss: 0.0156 - accuracy: 0.9952 - val_loss: 0.0302 - val_accuracy: 0.9903
Epoch 8/10
60000/60000 [=====] - 86s 1ms/step - loss: 0.0150 - accuracy: 0.9952 - val_loss: 0.0284 - val_accuracy: 0.9914
Epoch 9/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0124 - accuracy: 0.9960 - val_loss: 0.0290 - val_accuracy: 0.9914
Epoch 10/10
60000/60000 [=====] - 86s 1ms/step - loss: 0.0112 - accuracy: 0.9963 - val_loss: 0.0306 - val_accuracy: 0.9908
```

```
In [45]: score = model3.evaluate(X_test, y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```


10000/10000 [=====] - 4s 436us/step
Test loss: 0.03064428691338148
Test accuracy: 0.9908000230789185

```
In [46]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show()
```



Changing the dropout rate to 0.8

```
In [47]: model3=Sequential() # Initializing the model

# First ConvNet
model3.add(Conv2D(32,kernel_size=(3,3),
                  activation='relu',
                  input_shape=input_shape))
model3.add(MaxPooling2D(pool_size=(2,2)))

model3.add(Conv2D(64,kernel_size=(3,3),
                  activation='relu',padding='valid'))

model3.add(Dropout(0.8))

#hidden_layer
model3.add(Flatten())

model3.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model3.add(Dense(num_classes,activation='softmax'))

print(model3.summary())
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
=====		
conv2d_21 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_14 (MaxPooling)	(None, 13, 13, 32)	0
conv2d_22 (Conv2D)	(None, 11, 11, 64)	18496
dropout_20 (Dropout)	(None, 11, 11, 64)	0
flatten_8 (Flatten)	(None, 7744)	0
dense_15 (Dense)	(None, 256)	1982720
dense_16 (Dense)	(None, 10)	2570
=====		
Total params: 2,004,106		
Trainable params: 2,004,106		
Non-trainable params: 0		
None		

```
In [48]: model3.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [49]: # Training the model
```

```
history = model3.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/10
60000/60000 [=====] - 87s 1ms/step - loss: 0.2279 - accuracy: 0.9286 - val_loss: 0.0513 - val_accuracy: 0.9828
Epoch 2/10
60000/60000 [=====] - 89s 1ms/step - loss: 0.0831 - accuracy: 0.9742 - val_loss: 0.0337 - val_accuracy: 0.9890
Epoch 3/10
60000/60000 [=====] - 86s 1ms/step - loss: 0.0617 - accuracy: 0.9809 - val_loss: 0.0301 - val_accuracy: 0.9903
Epoch 4/10
60000/60000 [=====] - 86s 1ms/step - loss: 0.0514 - accuracy: 0.9841 - val_loss: 0.0268 - val_accuracy: 0.9912
Epoch 5/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0429 - accuracy: 0.9862 - val_loss: 0.0258 - val_accuracy: 0.9917
Epoch 6/10
60000/60000 [=====] - 86s 1ms/step - loss: 0.0387 - accuracy: 0.9870 - val_loss: 0.0231 - val_accuracy: 0.9915
Epoch 7/10
60000/60000 [=====] - 89s 1ms/step - loss: 0.0364 - accuracy: 0.9882 - val_loss: 0.0232 - val_accuracy: 0.9917
Epoch 8/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0327 - accuracy: 0.9895 - val_loss: 0.0216 - val_accuracy: 0.9922
Epoch 9/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0294 - accuracy: 0.9907 - val_loss: 0.0227 - val_accuracy: 0.9925
Epoch 10/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0271 - accuracy: 0.9911 - val_loss: 0.0204 - val_accuracy: 0.9919
```

```
In [50]: score = model3.evaluate(X_test, y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```


```
10000/10000 [=====] - 4s 432us/step
Test loss: 0.02036657058280398
Test accuracy: 0.9919000267982483
```

```
In [51]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show()
```



Changing dropout rate to 0.2


```
In [52]: model3=Sequential() # Initializing the model

# First ConvNet
model3.add(Conv2D(32,kernel_size=(3,3),
                  activation='relu',
                  input_shape=input_shape))
model3.add(MaxPooling2D(pool_size=(2,2)))

model3.add(Conv2D(64,kernel_size=(3,3),
                  activation='relu',padding='valid'))

model3.add(Dropout(0.2))

#hidden_layer
model3.add(Flatten())

model3.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))
model3.add(Dense(num_classes,activation='softmax'))

print(model3.summary())
```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d_23 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_15 (MaxPooling)	(None, 13, 13, 32)	0
conv2d_24 (Conv2D)	(None, 11, 11, 64)	18496
dropout_21 (Dropout)	(None, 11, 11, 64)	0
flatten_9 (Flatten)	(None, 7744)	0
dense_17 (Dense)	(None, 256)	1982720
dense_18 (Dense)	(None, 10)	2570
Total params: 2,004,106		
Trainable params: 2,004,106		
Non-trainable params: 0		

None

```
In [53]: model3.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

```
In [54]: # Training the model

history = model3.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 89s 1ms/step - loss: 0.1458 - accuracy: 0.9572 - val_loss: 0.0412 - val_accuracy:
0.9866
Epoch 2/10
60000/60000 [=====] - 87s 1ms/step - loss: 0.0413 - accuracy: 0.9873 - val_loss: 0.0330 - val_accuracy:
0.9887
Epoch 3/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0273 - accuracy: 0.9911 - val_loss: 0.0331 - val_accuracy:
0.9890
Epoch 4/10
60000/60000 [=====] - 86s 1ms/step - loss: 0.0202 - accuracy: 0.9932 - val_loss: 0.0347 - val_accuracy:
0.9885
Epoch 5/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0154 - accuracy: 0.9950 - val_loss: 0.0313 - val_accuracy:
0.9896
Epoch 6/10
60000/60000 [=====] - 88s 1ms/step - loss: 0.0126 - accuracy: 0.9958 - val_loss: 0.0299 - val_accuracy:
0.9907
Epoch 7/10
60000/60000 [=====] - 87s 1ms/step - loss: 0.0098 - accuracy: 0.9967 - val_loss: 0.0302 - val_accuracy:
0.9918
Epoch 8/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0066 - accuracy: 0.9977 - val_loss: 0.0356 - val_accuracy:
0.9908
Epoch 9/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0075 - accuracy: 0.9975 - val_loss: 0.0321 - val_accuracy:
0.9922
Epoch 10/10
60000/60000 [=====] - 85s 1ms/step - loss: 0.0062 - accuracy: 0.9981 - val_loss: 0.0352 - val_accuracy:
0.9909
```

```
In [55]: score = model3.evaluate(X_test, y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```


10000/10000 [=====] - 4s 426us/step
Test loss: 0.03524074672339475
Test accuracy: 0.9908999800682068

```
In [56]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show()
```



5 layer convolution layer with 5*5 kernel size

```
In [57]: model4=Sequential() # Initializing the model

# First ConvNet
model4.add(Conv2D(32,kernel_size=(5,5),
                  activation='relu',
                  input_shape=input_shape))
model4.add(MaxPooling2D(pool_size=(2,2)))

model4.add(Conv2D(64,kernel_size=(5,5),
                  activation='relu'))

model4.add(BatchNormalization())

model4.add(Conv2D(128,kernel_size=(5,5),
                  activation='relu'))
model4.add(Dropout(0.5))

model4.add(Flatten())

#hidden_layer
model4.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model4.add(Dense(512,activation='relu',kernel_initializer='random_uniform'))

model4.add(Dropout(0.5))

model4.add(Dense(num_classes,activation='softmax'))

print(model4.summary())
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
conv2d_25 (Conv2D)	(None, 24, 24, 32)	832
max_pooling2d_16 (MaxPooling)	(None, 12, 12, 32)	0
conv2d_26 (Conv2D)	(None, 8, 8, 64)	51264
batch_normalization_7 (Batch Normalization)	(None, 8, 8, 64)	256
conv2d_27 (Conv2D)	(None, 4, 4, 128)	204928
dropout_22 (Dropout)	(None, 4, 4, 128)	0
flatten_10 (Flatten)	(None, 2048)	0
dense_19 (Dense)	(None, 256)	524544
dense_20 (Dense)	(None, 512)	131584
dropout_23 (Dropout)	(None, 512)	0
dense_21 (Dense)	(None, 10)	5130
Total params: 918,538		
Trainable params: 918,410		
Non-trainable params: 128		
None		

```
In [58]: model4.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

In [59]: *# Training the model*

```
history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10
60000/60000 [=====] - 134s 2ms/step - loss: 0.1540 - accuracy: 0.9521 - val_loss: 0.0855 - val_accuracy: 0.9843

Epoch 2/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.0571 - accuracy: 0.9838 - val_loss: 0.0429 - val_accuracy: 0.9863

Epoch 3/10
60000/60000 [=====] - 134s 2ms/step - loss: 0.0420 - accuracy: 0.9885 - val_loss: 0.0632 - val_accuracy: 0.9838

Epoch 4/10
60000/60000 [=====] - 139s 2ms/step - loss: 0.0354 - accuracy: 0.9904 - val_loss: 0.0400 - val_accuracy: 0.9900

Epoch 5/10
60000/60000 [=====] - 142s 2ms/step - loss: 0.0348 - accuracy: 0.9910 - val_loss: 0.0307 - val_accuracy: 0.9916

Epoch 6/10
60000/60000 [=====] - 139s 2ms/step - loss: 0.0294 - accuracy: 0.9919 - val_loss: 0.0318 - val_accuracy: 0.9918

Epoch 7/10
60000/60000 [=====] - 140s 2ms/step - loss: 0.0286 - accuracy: 0.9926 - val_loss: 0.0478 - val_accuracy: 0.9879

Epoch 8/10
60000/60000 [=====] - 132s 2ms/step - loss: 0.0222 - accuracy: 0.9942 - val_loss: 0.0405 - val_accuracy: 0.9910

Epoch 9/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.0236 - accuracy: 0.9938 - val_loss: 0.0400 - val_accuracy: 0.9902

Epoch 10/10
60000/60000 [=====] - 133s 2ms/step - loss: 0.0228 - accuracy: 0.9942 - val_loss: 0.0442 - val_accuracy: 0.9907

In [60]: `score = model4.evaluate(X_test, y_test, verbose=1)`
`print('Test loss:', score[0])`
`print('Test accuracy:', score[1])`


```
10000/10000 [=====] - 7s 730us/step
Test loss: 0.044226658638985825
Test accuracy: 0.9907000064849854
```

In [61]: `fig,ax = plt.subplots(1,1)`
`ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')`

List of epoch numbers
`list_of_epoch = list(range(1,nb_epoch+1))`

`train_loss = history.history['loss']`
`val_loss = history.history['val_loss']`

`ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")`
`ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")`
`plt.legend()`
`plt.grid()`
`plt.show()`



Changing the dropout rate to 0.8

```
In [62]: model4=Sequential() # Initializing the model

# First ConvNet
model4.add(Conv2D(32,kernel_size=(5,5),
                  activation='relu',
                  input_shape=input_shape))
model4.add(MaxPooling2D(pool_size=(2,2)))

model4.add(Conv2D(64,kernel_size=(5,5),
                  activation='relu'))

model4.add(BatchNormalization())

model4.add(Conv2D(128,kernel_size=(5,5),
                  activation='relu'))
model4.add(Dropout(0.8))

model4.add(Flatten())

#hidden_layer
model4.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model4.add(Dense(512,activation='relu',kernel_initializer='random_uniform'))

model4.add(Dropout(0.8))

model4.add(Dense(num_classes,activation='softmax'))

print(model4.summary())
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
=====		
conv2d_28 (Conv2D)	(None, 24, 24, 32)	832
max_pooling2d_17 (MaxPooling)	(None, 12, 12, 32)	0
conv2d_29 (Conv2D)	(None, 8, 8, 64)	51264
batch_normalization_8 (Batch Normalization)	(None, 8, 8, 64)	256
conv2d_30 (Conv2D)	(None, 4, 4, 128)	204928
dropout_24 (Dropout)	(None, 4, 4, 128)	0
flatten_11 (Flatten)	(None, 2048)	0
dense_22 (Dense)	(None, 256)	524544
dense_23 (Dense)	(None, 512)	131584
dropout_25 (Dropout)	(None, 512)	0
dense_24 (Dense)	(None, 10)	5130
=====		
Total params: 918,538		
Trainable params: 918,410		
Non-trainable params: 128		
None		

```
In [63]: model4.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

In [64]: *# Training the model*

```
history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.3136 - accuracy: 0.9008 - val_loss: 0.0799 - val_accuracy: 0.9799

Epoch 2/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.0907 - accuracy: 0.9767 - val_loss: 0.0639 - val_accuracy: 0.9834

Epoch 3/10
60000/60000 [=====] - 134s 2ms/step - loss: 0.0745 - accuracy: 0.9816 - val_loss: 0.0393 - val_accuracy: 0.9904

Epoch 4/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.0632 - accuracy: 0.9840 - val_loss: 0.0421 - val_accuracy: 0.9906

Epoch 5/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.0656 - accuracy: 0.9847 - val_loss: 0.0316 - val_accuracy: 0.9923

Epoch 6/10
60000/60000 [=====] - 135s 2ms/step - loss: 0.0590 - accuracy: 0.9855 - val_loss: 0.0375 - val_accuracy: 0.9912

Epoch 7/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.0555 - accuracy: 0.9866 - val_loss: 0.0328 - val_accuracy: 0.9921

Epoch 8/10
60000/60000 [=====] - 129s 2ms/step - loss: 0.0494 - accuracy: 0.9890 - val_loss: 0.0301 - val_accuracy: 0.9922

Epoch 9/10
60000/60000 [=====] - 134s 2ms/step - loss: 0.0529 - accuracy: 0.9882 - val_loss: 0.0421 - val_accuracy: 0.9923

Epoch 10/10
60000/60000 [=====] - 131s 2ms/step - loss: 0.0485 - accuracy: 0.9888 - val_loss: 0.0411 - val_accuracy: 0.9920

In [65]:

```
score = model4.evaluate(X_test, y_test, verbose=1)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
10000/10000 [=====] - 8s 757us/step
Test loss: 0.04110135558335446
Test accuracy: 0.9919999837875366
```


In [66]:

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show()
```



Changing the dropout rate to 0.2

```
In [67]: model4=Sequential() # Initializing the model

# First ConvNet
model4.add(Conv2D(32,kernel_size=(5,5),
                  activation='relu',
                  input_shape=input_shape))
model4.add(MaxPooling2D(pool_size=(2,2)))

model4.add(Conv2D(64,kernel_size=(5,5),
                  activation='relu'))

model4.add(BatchNormalization())

model4.add(Conv2D(128,kernel_size=(5,5),
                  activation='relu'))
model4.add(Dropout(0.2))

model4.add(Flatten())

#hidden_layer
model4.add(Dense(256,activation='relu',kernel_initializer='random_uniform'))

model4.add(Dense(512,activation='relu',kernel_initializer='random_uniform'))

model4.add(Dropout(0.2))

model4.add(Dense(num_classes,activation='softmax'))

print(model4.summary())
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
conv2d_31 (Conv2D)	(None, 24, 24, 32)	832
max_pooling2d_18 (MaxPooling)	(None, 12, 12, 32)	0
conv2d_32 (Conv2D)	(None, 8, 8, 64)	51264
batch_normalization_9 (Batch Normalization)	(None, 8, 8, 64)	256
conv2d_33 (Conv2D)	(None, 4, 4, 128)	204928
dropout_26 (Dropout)	(None, 4, 4, 128)	0
flatten_12 (Flatten)	(None, 2048)	0
dense_25 (Dense)	(None, 256)	524544
dense_26 (Dense)	(None, 512)	131584
dropout_27 (Dropout)	(None, 512)	0
dense_27 (Dense)	(None, 10)	5130
Total params: 918,538		
Trainable params: 918,410		
Non-trainable params: 128		
None		

```
In [68]: model4.compile(optimizer='adam',
                       loss='categorical_crossentropy',
                       metrics=['accuracy'])
```

In [69]: *# Training the model*

```
history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/10
60000/60000 [=====] - 132s 2ms/step - loss: 0.1256 - accuracy: 0.9615 - val_loss: 0.0762 - val_accuracy: 0.9854
Epoch 2/10
60000/60000 [=====] - 134s 2ms/step - loss: 0.0473 - accuracy: 0.9866 - val_loss: 0.0387 - val_accuracy: 0.9881
Epoch 3/10
60000/60000 [=====] - 131s 2ms/step - loss: 0.0355 - accuracy: 0.9898 - val_loss: 0.0367 - val_accuracy: 0.9896
Epoch 4/10
60000/60000 [=====] - 132s 2ms/step - loss: 0.0297 - accuracy: 0.9916 - val_loss: 0.0579 - val_accuracy: 0.9829
Epoch 5/10
60000/60000 [=====] - 133s 2ms/step - loss: 0.0237 - accuracy: 0.9937 - val_loss: 0.0329 - val_accuracy: 0.9920
Epoch 6/10
60000/60000 [=====] - 132s 2ms/step - loss: 0.0207 - accuracy: 0.9938 - val_loss: 0.0552 - val_accuracy: 0.9866
Epoch 7/10
60000/60000 [=====] - 130s 2ms/step - loss: 0.0218 - accuracy: 0.9942 - val_loss: 0.0320 - val_accuracy: 0.9922
Epoch 8/10
60000/60000 [=====] - 133s 2ms/step - loss: 0.0178 - accuracy: 0.9950 - val_loss: 0.0344 - val_accuracy: 0.9919
Epoch 9/10
60000/60000 [=====] - 134s 2ms/step - loss: 0.0155 - accuracy: 0.9955 - val_loss: 0.0542 - val_accuracy: 0.9898
Epoch 10/10
60000/60000 [=====] - 134s 2ms/step - loss: 0.0172 - accuracy: 0.9953 - val_loss: 0.0381 - val_accuracy: 0.9927

In [70]: `score = model4.evaluate(X_test, y_test, verbose=1)`
`print('Test loss:', score[0])`
`print('Test accuracy:', score[1])`


10000/10000 [=====] - 7s 745us/step
Test loss: 0.038139941784558276
Test accuracy: 0.9926999807357788

In [71]: `fig,ax = plt.subplots(1,1)`
`ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')`

List of epoch numbers
`list_of_epoch = list(range(1,nb_epoch+1))`

`train_loss = history.history['loss']`
`val_loss = history.history['val_loss']`

`ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")`
`ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")`
`plt.legend()`
`plt.grid()`
`plt.show();`



Conclusions


```
In [73]: # Please compare all your models using Prettytable Library
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Numer of Layers", "Kernel Size", "BatchNormalization", "Dropout", "Accuracy"]

x.add_row(["3", '3*3', 'NO', '0.5', 99.08])
x.add_row(["3", '3*3', 'NO', '0.8', 99.19])
x.add_row(["3", '3*3', 'NO', '0.2', 99.08])

x.add_row(["4", '3*3', 'YES', '0.5', 99.29])
x.add_row(["4", '3*3', 'YES', '0.8', 93.98])
x.add_row(["4", '3*3', 'YES', '0.2', 99.21])

x.add_row(["4", '5*5', 'YES', '0.5', 99.43])
x.add_row(["4", '5*5', 'YES', '0.8', 98.33])
x.add_row(["4", '5*5', 'YES', '0.2', 99.29])

x.add_row(["5", '5*5', 'YES', '0.5', 99.07])
x.add_row(["5", '5*5', 'YES', '0.8', 99.19])
x.add_row(["5", '5*5', 'YES', '0.2', 99.26])

print(x)
```

Numer of Layers	Kernel Size	BatchNormalization	Dropout	Accuracy
3	3*3	NO	0.5	99.08
3	3*3	NO	0.8	99.19
3	3*3	NO	0.2	99.08
4	3*3	YES	0.5	99.29
4	3*3	YES	0.8	93.98
4	3*3	YES	0.2	99.21
4	5*5	YES	0.5	99.43
4	5*5	YES	0.8	98.33
4	5*5	YES	0.2	99.29
5	5*5	YES	0.5	99.07
5	5*5	YES	0.8	99.19
5	5*5	YES	0.2	99.26